

MYSTERY OF THE TRINITY

Audio-to-book production workflow and autonomous OpenCode handoff

Project: /Users/fr.peter/OpenCode/Mystery of the Trinity
Prepared: 16 September 2026
Author/source speaker: Fr. Peter Gruber, C.O.
Objective: Preserve seven recordings, create auditable transcripts, and develop a coherent book manuscript without inventing theology, quotations, or citations.

EXECUTIVE DECISION

Use a two-layer pipeline: (1) immutable, timestamped transcription records; then (2) separately versioned editorial and chapter files. Default to a local, word-timestamped Whisper implementation for privacy and resumability. Use a hosted transcription API only if speed is worth uploading the recordings and the owner explicitly accepts that transfer. Never draft book prose directly from audio without preserving the raw transcript and evidence trail.

Definition of done

- Every audio file has a checksum, machine transcript, human-reviewed clean transcript, review log, and completion marker.
- Every quotation, Scripture passage, doctrinal assertion, and source citation is either verified against an authoritative edition or visibly marked **[VERIFY]**.
- Six conference chapters and a separately assessed engineering-talk appendix/bonus chapter are assembled into one manuscript while retaining the author's cadence, pastoral examples, and theological emphases.
- No completed artifact is silently overwritten; every stage is deterministic, restartable, and logged.

Inventory at inspection

ID	Recording	Duration	Size	Initial role
C01	Icons of the Trinity	24:49	12.2 MB	Chapter 1
C02	Mystery	18:52	9.2 MB	Chapter 2
C03	Gift and Liturgy	23:58	11.8 MB	Chapter 3
C04	Relationship	13:26	6.6 MB	Chapter 4
C05	Intimacy	23:13	11.3 MB	Chapter 5
C06	Evangelization	15:52	7.8 MB	Chapter 6
E01	Engineering Mystery - Retreat Talk for Engineering Students	16:06	22.2 MB	Assess for appendix/bonus

Total audio: approximately 2:46:17. Existing companion file: "The Mystery of the Trinity Retreat 3-14.docx," containing the six-part outline, prayers, quotations, Scripture passages, canon-law excerpts, and Catechism material. Treat durations as container-reported estimates.

CONTENTS

1. Guardrails and recommended tools
 2. Canonical folder structure and names
 3. Exact end-to-end workflow
 4. Transcription specification
 5. Editorial rules and authorial voice
 6. Theological and citation verification
 7. Chapter architecture and synthesis
 8. Quality control and acceptance gates
 9. Resumability, idempotency, and logging
 10. OpenCode autonomous execution prompts
- Appendices: commands, manifests, checklists, and stop conditions

1. GUARDRAILS AND RECOMMENDED TOOLS

Non-negotiable operating rules

- **Preserve originals.** Audio and the source DOCX are read-only evidence. Never rename, normalize, trim, convert, or edit them in place.
- **Separate evidence from interpretation.** Machine transcript, reviewed transcript, editorial transcript, chapter draft, and source ledger are distinct artifacts.
- **No silent correction.** Ambiguous words, names, doctrine, quotations, and citations receive timestamps and flags; uncertainty is not smoothed away.
- **No fabricated support.** If a source cannot be verified, retain the author's meaning but mark the claim or wording for human review.
- **Human publication gate.** OpenCode may prepare a manuscript, but a named human approves doctrinal accuracy, quotation permissions, and final prose.

Preferred transcription stack

Need	Preferred	Fallback / note
Audio inspection	ffprobe / ffmpeg	macOS afinfo for duration
Local ASR	faster-whisper, large-v3, word timestamps	whisper.cpp or MLX Whisper on Apple Silicon
Hosted ASR	Current OpenAI transcription endpoint	Only with owner approval; pin model in manifest
Speaker handling	Single-speaker default + event tags	pyannote only if Q&A requires diarization
Text processing	Python 3 + JSON/Markdown	Avoid opaque binary-only output
Verification	Authoritative publisher/Church sources	Web sources saved with URL + access date
Versioning	Git for text/scripts/manifests	Git LFS only if explicitly configured; do not add originals by default

Model policy

Record the exact engine, model identifier, package version, decoding settings, language, timestamp granularity, and execution date. Use one pinned configuration for the whole batch. A later model upgrade creates a new run ID; it does not overwrite the first run.

2. CANONICAL FOLDER STRUCTURE AND NAMES

```
Mystery of the Trinity/  
  audio_original/      # copies of the 7 source M4A files; immutable  
  source_material/     # source DOCX and extracted text  
  manifests/           # inventory.csv, checksums.sha256, run manifests  
  scripts/             # deterministic pipeline and validators  
  work/  
    normalized_audio/  # optional WAV/FLAC derivatives  
    chunks/           # generated chunks + chunk manifests  
    asr/               # raw engine JSON by run ID  
    review/            # issue queues and listening notes  
  transcripts/  
    raw/               # timestamped machine transcript Markdown  
    clean/             # human-reviewed, faithful transcript  
    editorial/         # lightly edited oral-to-written version  
  research/  
    source_ledger.csv  
    verified_excerpts/  
    unresolved.md  
  chapters/  
    outlines/  
    drafts/  
    reviewed/  
  manuscript/  
    Mystery_of_the_Trinity.md  
    notes.md  
  logs/  
  archive/
```

Stable IDs and filenames

- C01_icons_of_the_trinity, C02_mystery, C03_gift_and_liturgy, C04_relationship, C05_intimacy, C06_evangelization, E01_engineering_mystery.
- Derivative pattern: **{id}.{stage}.{run_id}.{ext}**, for example **C01.asr.20260916_local_whisper.json**. The approved clean transcript may also have a stable pointer file **C01.clean.md** containing provenance.
- Chunk pattern: **{id}__000000-001500.wav** using zero-padded start/end seconds. Do not call chunks “part1”; time ranges make retries and audits deterministic.
- Run ID pattern: **YYYYMMDD_engine_model_short**. Do not use “final,” “final2,” or ambiguous dates.

Manifest minimum fields

```
recording_id, source_path, basename, bytes, sha256, duration_seconds,  
codec, sample_rate, channels, source_mtime, status, asr_run_id,  
transcript_path, reviewer, reviewed_at, notes
```

3. EXACT END-TO-END WORKFLOW

- 1 **Initialize safely.** Create the structure; copy originals into audio_original only if absent; mark originals read-only; create .gitignore entries for large audio and generated chunks unless the owner wants them versioned.
- 2 **Inventory and fingerprint.** Generate inventory.csv and SHA-256 checksums. Compare the seven expected IDs to discovered files. Stop on duplicates, zero-byte files, unreadable media, or checksum drift.
- 3 **Extract source material.** Convert the DOCX to plain text/Markdown while retaining the DOCX unchanged. Create a preliminary source ledger from its named quotations, passages, and numbered sources.
- 4 **Preflight audio.** Record duration/codecs/sample rate/channels. Listen to the first, middle, and last 30 seconds of each file. Log clipping, silence, music, audience questions, and speaker changes.
- 5 **Normalize only if needed.** ASR can usually accept M4A directly. If conversion is necessary, create mono 16 kHz PCM WAV derivatives; never alter originals.
- 6 **Transcribe batch.** Process stable IDs in order with pinned settings, word timestamps, JSON output, and per-file logs. Write to a temporary file, validate it, then atomically rename it and write a .done marker.
- 7 **Review transcript.** Correct against audio at normal speed, then spot-check difficult passages at 0.8x. Resolve proper nouns and obvious ASR errors; retain uncertainty flags with timestamps.
- 8 **Create clean and editorial layers.** Clean transcript removes disfluencies under the rules below. Editorial transcript converts speech to readable prose without adding claims or changing theology.
- 9 **Verify sources.** Populate source_ledger.csv. Compare every quotation and numbered citation with an authoritative source. Record exact wording, edition, locator, URL/ISBN, access date, and disposition.
- 10 **Segment and synthesize.** Build a timestamp-linked outline per talk, map themes across talks, decide where E01 contributes unique material, and draft one chapter at a time from approved transcript excerpts.
- 11 **Run chapter QA.** Require traceability, voice preservation, citation closure, cross-chapter consistency, and human theological/editorial review before promoting drafts.
- 12 **Assemble manuscript.** Combine only approved chapter versions; generate a global bibliography/notes list; run final terminology, cross-reference, and completeness checks.

Recommended order: pilot C01 end to end first. Freeze conventions after that pilot, then batch C02-C06 and E01. This prevents seven files from inheriting an early formatting or editorial mistake.

4. TRANSCRIPTION SPECIFICATION

Batch strategy

- First attempt: one job per whole recording. These files are short enough (13-46 minutes) that whole-file context is valuable.
- If memory, timeout, or API limits require chunking, use 15-minute chunks with 20-second overlap. Keep chunk boundaries deterministic and record them in a manifest.
- When merging overlaps, match timestamped words; never concatenate blindly. Preserve the earlier segment until the semantic boundary, then append only non-duplicate words.
- Limit parallel jobs to available CPU/GPU/memory. Start with one; raise concurrency only after measuring one complete file. Determinism and recoverability matter more than peak throughput.
- Force or strongly bias English only after confirming all recordings are English. Supply a short vocabulary hint where supported: Trinity, Trinitarian, Eucharist, liturgy, Ratzinger, Rublev, perichoresis, consubstantial, evangelization.

Speaker and timestamp policy

Situation	Required rendering
Main speaker	SPEAKER_01 in raw JSON; display as "Fr. Peter" only after identity confirmation.
Audience / question	AUDIENCE_01, AUDIENCE_02 when distinguishable; otherwise [Audience member].
Unclear speaker	[Unknown speaker]; do not guess.
Non-speech	[laughter], [applause], [pause], [inaudible 00:12:34], only when editorially meaningful.
Raw timestamps	Segment start for every paragraph; word timestamps remain in JSON.
Clean timestamps	At least every 60-90 seconds and at topic/speaker changes.
Book draft	No timestamps in prose; retain source anchors in hidden/editor notes.

Raw transcript header

```
---
recording_id: C01
source_sha256: <64 hex chars>
asr_run_id: 20260916_local_whisper_large-v3
engine: faster-whisper
model: large-v3
language: en
word_timestamps: true
status: machine_raw
---
```

5. EDITORIAL RULES AND AUTHORIAL VOICE

Three transcript layers

Layer	May do	Must not do
Raw	Add timestamps/speaker labels; render engine output	Correct, summarize, or polish silently
Clean	Remove false starts, repeated filler, obvious stutters; fix ASR errors; remove filler	Rephrase or replace characteristic phrasing; add theology
Editorial	Repair oral syntax, collapse repetition, add headings, move a passage with purpose	Insert connectives, transitions, citations, or certainty
Chapter	Synthesize approved material into book prose with source and chapter change log	Present generated prose as an author quotation or erase unresolved doubts

Verbatim-to-clean rules

- Delete “um,” “uh,” abandoned syllables, mic checks, and duplicated words unless they express hesitation that matters.
- Keep meaningful repetition, rhetorical questions, pastoral asides, humor, prayers, and direct address when they carry the author’s voice.
- Convert spoken punctuation and obvious contractions consistently, but do not modernize theological terms or “improve” doctrinal formulations.
- Mark uncertain audio as **[unclear 00:00:00]**. A plausible guess may be added as **[possibly: term]**, never substituted silently.
- Use **[editorial insertion]** only in transcripts. Book prose may have transitions, but the chapter change log must identify generated connective material.
- Quoted wording heard in audio stays in quotation marks only if verified or labeled **[quotation wording unverified]**.

Voice profile to build before chapter drafting

- Create research/voice_profile.md using only reviewed transcripts. Capture characteristic sentence length, preferred theological vocabulary, use of first/second person, humor, analogy, prayerful register, and typical movement from doctrine to lived application.
- Collect 15-25 short, timestamped voice exemplars across all six conferences. Do not turn the profile into caricature; it is a restraint on generic AI prose.
- For every drafted section, run a “voice delta” check: identify vocabulary or claims absent from the source, overly academic smoothing, and lost pastoral directness.

6. THEOLOGICAL AND CITATION VERIFICATION

Source hierarchy

- 1 **Sacred Scripture:** use the edition selected by the author/publisher. Verify book, chapter, verse range, wording, and translation attribution. Do not silently mix translations.
- 2 **Magisterial and canonical texts:** use official Vatican or authorized episcopal/publisher sources where possible. Verify paragraph/canon numbers and edition date.
- 3 **Patristic, saint, papal, and theological works:** verify against a named edition or official archive; capture translator, publication data, and page/section locator.
- 4 **Secondary commentary:** use reputable scholarly editions. It may explain a claim but should not replace verification of a primary quotation.
- 5 **Unverified web reproductions:** use only as leads. Never cite quotation aggregators as final authority.

Source ledger schema

```
source_id, chapter_id, transcript_anchor, item_type, attributed_author,
work_title, quoted_or_paraphrased_text, claimed_locator, verified_text,
edition_or_translation, publisher, year, page_or_section, canonical_url,
accessed_at, verification_status, discrepancy, permission_status, reviewer
```

Status vocabulary

- **VERIFIED_EXACT:** wording and locator match the selected edition.
- **VERIFIED_MINOR_VARIANT:** differences are documented (translation, punctuation, edition).
- **PARAPHRASE_CONFIRMED:** idea is supported, but wording is not a direct quotation.
- **LOCATOR_MISSING:** work/source found; exact place not yet established.
- **ATtribution_UNCONFIRMED:** wording or author cannot be established.
- **PERMISSION_REVIEW:** accurate but length/rights require publisher review.
- **DOCTRINAL_REVIEW:** faithful transcription but theological precision needs a qualified human.

Verification procedure

- 1 Extract every quoted phrase, citation-like mention, Scripture reference, Church document, canon, saint, artwork, historical claim, Greek/Latin term, and attributed anecdote.
- 2 Link the item to recording ID and timestamp, then search the existing DOCX first. Its content is evidence of intent, not proof of accuracy.
- 3 Locate an authoritative source; record bibliographic details and a short verified excerpt. For web sources, save canonical URL and access date.
- 4 Compare exact wording and locator. Record discrepancies rather than normalizing them invisibly.
- 5 Decide: quote exactly, quote with a stated translation, paraphrase with citation, omit, or escalate.
- 6 A qualified theological reviewer closes DOCTRINAL_REVIEW items. A publishing/rights reviewer closes PERMISSION_REVIEW items.

The companion DOCX already names Ratzinger, Pope Benedict XVI, the Code of Canon Law, the Catechism, St. John Henry Newman, St. Elizabeth of the Trinity, St. Patrick's Breastplate, Rublev's Trinity, and multiple Scripture passages. Seed the ledger with these before transcript review so ASR variants can be matched consistently.

7. CHAPTER ARCHITECTURE AND SYNTHESIS

ID	Working chapter	Primary task
C01	Icons of the Trinity	Establish images, revelation, and entry points into the mystery.
C02	Mystery	Clarify mystery as inexhaustible reality, not mere puzzle.
C03	Gift and Liturgy	Develop divine self-gift and liturgical participation.
C04	Relationship	Show Trinitarian communion as criterion for human relation.
C05	Intimacy	Move toward prayer, indwelling, desire, and transformation.
C06	Evangelization	Carry contemplation outward into witness and mission.
E01	Engineering Mystery	Use only unique, complementary material; appendix, interlude, or source pool after comparison.

Segmentation process

- 1 Create a talk map with segments of roughly 2-7 minutes based on idea changes, not arbitrary length. Each row: start/end, descriptive heading, one-sentence claim, evidence/examples, cited sources, memorable wording, and flags.
- 2 Create a cross-talk theme matrix: rows are themes; columns are C01-C06 and E01; cells contain timestamps. Identify repetitions, developments, tensions, and unique material.
- 3 Write a one-page chapter brief: reader promise, controlling question, core claim, emotional/spiritual movement, essential sources, opening image, closing invitation, and excluded material.
- 4 Draft a detailed outline with transcript anchors before prose. Every major claim needs at least one source anchor or a clearly labeled editorial bridge.
- 5 Draft in small sections. After each section, append an origin note listing recording/timestamp spans, verified sources, and any new connective prose.
- 6 Compare the draft to the clean transcript and voice profile. Restore distinctive examples and cadence where polishing has made the prose generic.
- 7 Move surplus but valuable material into a parking-lot file; do not force all audio into the book.

Recommended chapter pattern

- Opening concrete image, question, or experience from the talk.
- The theological claim stated accessibly and with appropriate humility before mystery.
- Scriptural/Church/theological witness, verified and cited.
- Pastoral or human implication; preserve the author's characteristic movement between contemplation and life.
- Return to the opening motif and invite prayer, reflection, liturgy, relationship, or mission as appropriate.

8. QUALITY CONTROL AND ACCEPTANCE GATES

Gate	Checks	Pass condition
G0 Intake	Seven audio files + DOCX; readable; checksum manifest	No missing/duplicate/corrupt source; immutable copy established
G1 ASR	JSON parses; timestamps monotonic; duration coverage; no [unclear]	At least 99% timeline coverage excluding legitimate silence; all anomalies logged
G2 Transcript	Full listen; names/terms; speaker labels; uncertainty flags	Reviewer signs clean transcript; unresolved items queued
G3 Sources	Ledger completeness; wording; locators; translation; rights	No unmarked citation claim; blockers have owners/status
G4 Chapter	Traceability; argument; voice; repetition; factual/theological relevance	Material claims anchored; reviewer approval recorded
G5 Manuscript	Order; cross-references; glossary/style; bibliography; complete	All approved chapters assembled; no [VERIFY]/[TODO] unless explicitly accepted

Required spot checks

- For every file: inspect first 2 minutes, last 2 minutes, and five randomly selected 60-second windows; additionally inspect every [unclear], quotation, proper noun, and speaker transition.
- Calculate and report unresolved flags per 1,000 words. Do not invent a “confidence percentage” unless it is tied to a defined sampling method.
- Compare transcript end timestamp with media duration; flag gaps over 10 seconds that are not documented silence/music.
- Search the repository before promotion for [VERIFY], [TODO], [unclear], `ATtribution_UNCONFIRMED`, `LOCATOR_MISSING`, and `DOCTRINAL_REVIEW`.
- Run a chapter-source audit: sample at least ten substantive sentences per chapter and trace each to transcript anchors or verified sources.

Human sign-offs

Assign names, not merely “reviewed”: transcript reviewer; author/voice reviewer; theological reviewer; citation/rights reviewer; final editor. One person may hold multiple roles, but each sign-off must include date, artifact hash, decision, and remaining exceptions.

9. RESUMABILITY, IDEMPOTENCY, AND LOGGING

Rules that prevent lost work

- Every stage computes an input fingerprint from source SHA-256 plus configuration hash. The same fingerprint must produce the same target filename or be skipped.
- Write outputs as **.partial** in the destination directory, validate, fsync/close, then atomically rename. A crash may leave .partial files; never treat them as complete.
- A **.done.json** marker records input hashes, output hashes, command/config, start/end time, tool versions, and validation result. Existence alone is insufficient; hashes must match.
- If a valid output exists, skip it by default. Reprocessing requires **--force-new-run** and a new run ID. Never overwrite a reviewed transcript or approved chapter.
- Edits to clean/editorial/chapter files create versioned commits and update provenance. Raw ASR output is immutable after validation.
- Failures are recorded per recording in logs/run_id.jsonl; the batch continues where safe and exits nonzero if any required file failed.
- On restart, reconcile manifest, outputs, .partial files, and .done markers. Resume missing work only; quarantine inconsistent artifacts for inspection.

State machine

```
DISCOVERED -> INVENTORIED -> PREFLIGHT_OK -> ASR_DONE -> TRANSCRIPT_REVIEWED  
-> SOURCES_VERIFIED -> OUTLINED -> DRAFTED -> CHAPTER_APPROVED -> ASSEMBLED
```

Any state may also be BLOCKED with: reason, owner, created_at, next_action.
Only an explicit review record promotes human-gated states.

Suggested safe shell patterns

```
# Inventory without modifying sources  
find audio_original -type f -name '*.m4a' -print0 | sort -z  
shasum -a 256 audio_original/*.m4a > manifests/checksums.sha256  
  
# Optional derivative; create only if missing and verify before rename  
ffmpeg -nostdin -v error -i SOURCE.m4a -ac 1 -ar 16000 -c:a pcm_s16le OUT.partial.wav  
# validate OUT.partial.wav, then: mv OUT.partial.wav OUT.wav  
  
# Search for unresolved editorial markers before manuscript promotion  
rg -n '\[(VERIFY|TODO|unclear)|ATtribution_UNCONFIRMED|DOCTRINAL_REVIEW' \  
  transcripts chapters manuscript research
```

Do not run destructive cleanup automatically. Generated chunks may be removed only after their parent transcript is validated and their regeneration recipe is recorded; originals and reviewed text are never cleanup targets.

10. OPENCODE AUTONOMOUS EXECUTION PROMPTS

Master operating prompt

You are operating the Mystery of the Trinity audio-to-book pipeline.
Treat `audio_original/` and `source_material/` as immutable evidence.
Read this handoff and manifests before acting. Work one declared stage at a time.
Before writing, state inputs, expected outputs, validation, and skip criteria.
Use stable recording IDs. Never overwrite a valid artifact; create a new run/version.
Preserve raw ASR JSON and timestamped transcript before editing.
Do not guess inaudible words, speakers, quotations, citations, or doctrine.
Use `[unclear HH:MM:SS]` and `source-ledger` statuses.
Every drafted claim must trace to transcript timestamps, a verified source, or a clearly labeled editorial bridge. Stop at human gates and produce a review queue.
At the end, report created/skipped/failed/blocked items and exact next action.

Prompt A - initialize and inventory

Inspect the project without modifying original files. Create the canonical folder structure, stable ID mapping, `inventory.csv`, `checksums.sha256`, and an intake report. Confirm exactly seven recordings and the companion DOCX. Record duration, codec, sample rate, channels, byte size, mtime, and SHA-256. Extract the DOCX to a derivative Markdown file. Stop on source ambiguity or unreadable media. Do not transcribe yet.

Prompt B - pilot transcription

Run the pinned transcription workflow on C01 only. Prefer whole-file local ASR with English and word timestamps. Save engine-native JSON, timestamped raw Markdown, logs, configuration, and a validated `.done.json` marker. Do not clean the text. Report audio anomalies, likely speaker count, uncertain terminology, timeline coverage, and a proposal for transcript formatting. Await approval before batching.

Prompt C - batch transcription

Using the approved C01 configuration and format, process C02-C06 and E01. Skip outputs whose input/config/output hashes validate. Resume incomplete work only. For each recording produce raw JSON, raw Markdown, validation report, log, and completion marker. Never overwrite C01 or any reviewed artifact. Return a batch summary with failures and retry commands.

Prompt D - transcript review package

For each raw transcript, create a review copy and issue list. Correct only against the audio. Apply the verbatim-to-clean rules in this handoff. Preserve timestamps at topic and speaker boundaries. Flag uncertainty, quotations, names, doctrine, and citations. Produce a change summary and request human transcript sign-off; do not promote to clean/ without that sign-off.

AUTONOMOUS PROMPTS, CONTINUED

Prompt E - source ledger and verification

Seed source_ledger.csv from the source DOCX, then extract candidate sources from all reviewed transcripts. Link every item to recording ID and timestamp. Verify against authoritative editions in the stated hierarchy. Record exact bibliographic data, discrepancy, status, and permission/doctrinal flags. Never fill missing data from memory. Produce unresolved.md grouped by the human decision required.

Prompt F - chapter map

Using approved clean transcripts and the verified source ledger, create a timestamp-linked segment map for every recording and a cross-talk theme matrix. Propose six core chapter briefs following C01-C06. Compare E01 against the six talks and recommend appendix, interlude, source-pool, or omission with evidence. Do not draft prose yet. Surface duplicated material and unresolved structural choices.

Prompt G - chapter drafting

Draft only the approved chapter ID from its approved brief, clean transcript, voice profile, and verified source ledger. Preserve characteristic examples and pastoral cadence. Add no unsupported theology or citations. Maintain a sidecar provenance file mapping each section to timestamps/sources and identifying editorial bridges. Run the chapter QA checklist, then stop for voice and theological review.

Prompt H - manuscript assembly

Assemble only approved chapter versions into manuscript/Mystery_of_the_Trinity.md. Do not rewrite approved prose during assembly. Normalize headings, notes, bibliography, terminology, and cross-references. Scan all project text for unresolved markers. Generate a completion report listing included artifact hashes, approvals, exceptions, and blockers. Do not label the manuscript final while any blocking marker remains.

Required report format after every run

```
RUN ID / STAGE / START-END
INPUTS + HASHES
CREATED
SKIPPED (with reason)
FAILED (with exact error)
BLOCKED (decision owner + next action)
VALIDATION RESULTS
UNRESOLVED FLAGS
SAFE NEXT COMMAND / HUMAN GATE
```

APPENDIX A - IMPLEMENTATION BLUEPRINT

Suggested script interfaces

```
scripts/01_inventory.py --project .
scripts/02_extract_docx.py --input source_material/...docx
scripts/03_preflight_audio.py --manifest manifests/inventory.csv
scripts/04_transcribe.py --id C01 --config manifests/asr_config.yaml
scripts/05_render_transcript.py --id C01 --run-id RUN_ID
scripts/06_validate_asr.py --id C01 --run-id RUN_ID
scripts/07_build_review_queue.py --id C01
scripts/08_extract_sources.py --all-reviewed
scripts/09_validate_repository.py --gate G3
scripts/10_assemble_manuscript.py --approved-only
```

Example ASR configuration

```
engine: faster-whisper
model: large-v3
device: auto
compute_type: auto
language: en
word_timestamps: true
vad_filter: true
beam_size: 5
temperature: 0
chunk_seconds: null
chunk_overlap_seconds: 20
vocabulary_hint:
  - Trinity
  - Trinitarian
  - Eucharist
  - Ratzinger
  - Rublev
  - perichoresis
  - consubstantial
```

Validation pseudocode

```
for recording in manifest:
    fingerprint = sha256(source_sha256 + canonical_json(config))
    if valid_done_marker(recording, fingerprint):
        skip(recording)
        continue
    output_tmp = destination + '.partial'
    transcribe(recording.source, output_tmp, config)
    assert json_pares(output_tmp)
    assert timestamps_are_monotonic(output_tmp)
    assert duration_coverage_is_acceptable(output_tmp, recording.duration)
    atomic_rename(output_tmp, destination)
    write_done_marker_atomically(destination, fingerprint, sha256(destination))
```

Version-control boundaries

- Commit: manifests, scripts, logs without secrets, Markdown transcripts, source ledger, chapter files, review records, and configuration.
- Usually ignore: audio_original copies, normalized audio, chunks, model caches, temporary files, API response caches containing sensitive data.
- Never commit: API keys, credentials, tokens, private reviewer contact data, or copyrighted source scans without permission.
- If originals remain only in the top-level legacy location, document that path in the manifest and back them up before any reorganization.

APPENDIX B - REVIEW CHECKLISTS

Transcript reviewer

- Source hash matches manifest; correct recording ID and ASR run recorded.
- Beginning and ending are complete; long gaps are explained.
- Speaker labels are evidence-based; audience speech is not attributed to the author.
- Proper nouns and theological terms checked against audio/source map.
- Quotations and citations flagged; uncertainty retains timestamps.
- Clean edits obey stated rules; meaningful rhetoric and prayer remain.
- Reviewer name/date/artifact hash recorded.

Chapter reviewer

- Chapter has a clear promise and progression, not merely a transcript summary.
- Every substantive theological/factual claim traces to audio or verified source.
- Direct quotations match the selected edition and locator; paraphrases are not quoted.
- Author's vocabulary, examples, humility, and pastoral cadence remain recognizable.
- Generated bridges are accurate, minimal, and disclosed in provenance notes.
- Repetition with other chapters is purposeful; contradictions or terminology drift are resolved.
- Opening and conclusion serve the book's arc; surplus material is parked rather than forced.
- Voice and theological reviewers approve the hashed version.

Final manuscript release

- All expected chapters present in approved order; E01 disposition documented.
- No blocking markers remain; accepted exceptions are listed explicitly.
- Scripture translation and citation style are consistent.
- Bibliography/notes reconcile with source ledger; artwork and long quotations have rights review.
- Table of contents, headings, cross-references, and chapter titles agree.
- Final manuscript hash, source commit, approvals, and release date recorded.
- Originals, transcripts, ledger, and manuscript are backed up in at least two locations.

Stop conditions for OpenCode: missing or changed source file; conflicting IDs; corrupt media; unknown speaker whose identity affects attribution; inaudible wording that changes doctrine; unavailable primary source for a direct quotation; unresolved Scripture translation; possible rights issue; instruction to overwrite a reviewed artifact; or any request to infer an authorial claim not supported by the evidence.

APPENDIX C - PROJECT-SPECIFIC STARTING NOTES

Observed source alignment

- The six conference filenames match the six headings in the DOCX exactly enough to use C01-C06 as the initial chapter map.
- The engineering recording is substantially longer than any conference and is not listed in the DOCX's six-conference outline; therefore it should not automatically become Chapter 7. Compare it after transcript mapping.
- The DOCX supplies a strong preliminary authority list and many exact passages, but it also contains typographic and transcription-like irregularities. All wording and locators still require edition-level verification.
- Audio is low-bitrate spoken-word M4A (about 63-64 kbps). Preserve the originals; transcoding cannot restore lost detail and should be used only for ASR compatibility.

First autonomous session - exact scope

- 1 Create folders and a non-destructive inventory; do not move/delete originals.
- 2 Assign the seven stable IDs in this document.
- 3 Generate checksums, technical metadata, and extracted DOCX Markdown.
- 4 Create a source-ledger seed list from the DOCX.
- 5 Preflight-listen to each file and record speaker/audio observations.
- 6 Propose and pin the ASR configuration.
- 7 Transcribe C01 only and stop at the pilot review gate.

Ownership decisions to capture before publication

- Preferred Bible translation and whether quoted liturgical/prayer texts follow a particular edition.
- Book audience, desired chapter length, citation style, and whether the voice remains explicitly retreat/conference-like.
- Whether E01 is intended for the book, a bonus essay, or background only.
- Names of transcript, theological, rights, voice, and final editorial reviewers.
- Privacy decision for local versus hosted transcription and the backup location for originals.
- Publisher policy for permissions, AI-assisted editing disclosure, and archival retention.

Handoff complete. The immediate safe next action is inventory + C01 pilot, not full-book drafting.